1) How a GEMM code written in Triton gets converted to Nvidia PTX code using MLIR?
 Show code fragements after each lowering for a matrix multiplication of two matrices A and B of size 1023x1024.

2) Differentiate between the data types FP64, TF32, FP16, BF16, FP8. What is purpose data types with low precision?

3) How warps schedule tiled matrix multiplication on tensor cores?

3) Compare the Nvidia GPUs : A100, H100 and B100 GPU based on hardware manual